

# Augment and Reduce: Stochastic Inference for Large Categorical Distributions

Francisco J. R. Ruiz


April 27th, 2018



# Joint Work With



# Categorical Distributions



- ▶ A probability distribution on a set of  $K$  outcomes
- ▶ Normalized,  $\sum_k p_k = 1$
- ▶ Ubiquitous in machine learning and many other disciplines

## Our Contribution



- ▶ Goal: Speed up training for models with large categoricals ( $K \gg 1$ )
- ▶ Contribution: A fast algorithm with controlled complexity
- ▶ Key ideas: Variable augmentation, stochastic variational inference

# Softmax

- ▶ One widely applied parameterization of a categorical,

$$p(y = k \mid \psi) = \frac{e^{\psi_k}}{\sum_{k'} e^{\psi_{k'}}}$$

- ▶ Transforms reals into probabilities



## A Case Study: Classification

- ▶ Observations are features and labels,  $\{x_n, y_n\}_{n=1}^N$
- ▶ Each label  $y_n \in \{1, \dots, K\}$

## A Case Study: Classification

- ▶ Observations are features and labels,  $\{x_n, y_n\}_{n=1}^N$
- ▶ Each label  $y_n \in \{1, \dots, K\}$
- ▶ Each observation  $n$  is assigned a real value,

$$\psi_k^{(n)} = w_k^\top x_n$$

- ▶ Goal: Find the weights  $w = (w_1, \dots, w_K)$

## A Case Study: Classification

- ▶ Maximize the likelihood of the data w.r.t. the weights,

$$\text{find } \mathbf{w} \text{ to maximize } \mathcal{L}_{\text{log-lik}} = \sum_n \log p(y_n | x_n, \mathbf{w})$$

## A Case Study: Classification

- ▶ Maximize the likelihood of the data w.r.t. the weights,


$$\text{find } w \text{ to maximize } \mathcal{L}_{\text{log-lik}} = \sum_n \log p(y_n | x_n, w)$$

- ▶ Assuming the softmax transformation,

$$\log p(y_n | x_n, w) = \log \left( \frac{e^{w_{y_n}^\top x_n}}{\sum_{k'} e^{w_{k'}^\top x_n}} \right)$$

# A Case Study: Classification

- ▶ Optimization w.r.t.  $w$
- ▶ Gradient ascent



- ▶ The gradient is

$$\nabla_w \mathcal{L}_{\text{log-lik}} = \sum_n \nabla_w \log p(y_n | x_n, w)$$

## A Case Study: Classification


$$\nabla_w \log p(y_n | x_n, w) = \nabla_w \log \left( \frac{e^{w_{y_n}^\top x_n}}{\sum_{k'} e^{w_{k'}^\top x_n}} \right)$$

- ▶ **Problem:** Evaluating the gradient is  $\mathcal{O}(K)$

# A Case Study: Classification

$$\nabla_w \log p(y_n | x_n, w) = \nabla_w \log \left( \frac{e^{w_{y_n}^\top x_n}}{\sum_{k'} e^{w_{k'}^\top x_n}} \right)$$


- ▶ **Problem:** Evaluating the gradient is  $\mathcal{O}(K)$
- ▶ Evaluation is needed for each  $n = 1, \dots, N$  and at each iteration of gradient ascent



# A Case Study: Classification

$$\nabla_w \log p(y_n | x_n, w) = \nabla_w \log \left( \frac{e^{w_{y_n}^\top x_n}}{\sum_{k'} e^{w_{k'}^\top x_n}} \right)$$

- ▶ **Problem:** Evaluating the gradient is  $\mathcal{O}(K)$
- ▶ Evaluation is needed for each  $n = 1, \dots, N$  and at each iteration of gradient ascent



- ▶ For large values of  $K$ , this is prohibitive

## Large Categoricals

- ▶ The  $\mathcal{O}(K)$  cost is not unique to the softmax
- ▶ Other models (multinomial probit/logistic) are also  $\mathcal{O}(K)$

## Large Categoricals

- ▶ The  $\mathcal{O}(K)$  cost is not unique to the softmax
- ▶ Other models (multinomial probit/logistic) are also  $\mathcal{O}(K)$
- ▶ When  $K$  is large, this is not OK

## Large Categoricals

- ▶ The  $\mathcal{O}(K)$  cost is not unique to the softmax
- ▶ Other models (multinomial probit/logistic) are also  $\mathcal{O}(K)$
- ▶ When  $K$  is large, this is not OK
- ▶ Examples: language models, recommendation systems, discrete choice models, reinforcement learning

## Our Contribution

- ▶ An algorithm with reduced complexity,  $\mathcal{O}(|\mathcal{S}|)$  instead of  $\mathcal{O}(K)$
- ▶ Complexity controlled by parameter  $|\mathcal{S}|$

## Our Contribution

- ▶ An algorithm with reduced complexity,  $\mathcal{O}(|\mathcal{S}|)$  instead of  $\mathcal{O}(K)$
- ▶ Complexity controlled by parameter  $|\mathcal{S}|$
- ▶ Two steps
  1. Augment the model with an auxiliary variable
  2. Reduce complexity via subsampling (stochastic optimization)

## Let's Take A Step Back...

- ▶ Where does the softmax come from?

$$p(y = k \mid \psi) = \frac{e^{\psi_k}}{\sum_{k'} e^{\psi_{k'}}}$$

## The Utility Perspective

- ▶ Draw random errors i.i.d.,  $\varepsilon_k \sim \phi(\cdot)$

## The Utility Perspective

- ▶ Draw random errors i.i.d.,  $\varepsilon_k \sim \phi(\cdot)$
- ▶ Define a *utility* for each outcome  $k$ ,

$$\psi_k + \varepsilon_k$$

(mean utility plus noise)

## The Utility Perspective

- ▶ Draw random errors i.i.d.,  $\varepsilon_k \sim \phi(\cdot)$
- ▶ Define a *utility* for each outcome  $k$ ,

$$\psi_k + \varepsilon_k$$

(mean utility plus noise)

- ▶ Choose the outcome with the largest utility,

$$y = \arg \max_k (\psi_k + \varepsilon_k)$$

## The Utility Perspective

- ▶ Draw random errors i.i.d.,  $\varepsilon_k \sim \phi(\cdot)$
- ▶ Define a *utility* for each outcome  $k$ ,

$$\psi_k + \varepsilon_k$$

(mean utility plus noise)

- ▶ Choose the outcome with the largest utility,

$$y = \arg \max_k (\psi_k + \varepsilon_k)$$

- ▶ Integrate out the error terms ( $\varepsilon_k$ 's) to find the marginal  $p(y | \psi)$

## The Utility Perspective

- ▶ Different priors  $\phi(\varepsilon)$  lead to different categoricals

## The Utility Perspective

- ▶ Different priors  $\phi(\varepsilon)$  lead to different categoricals
- ▶ For  $\phi(\varepsilon) = \text{Gumbel}(\varepsilon \mid 0, 1)$ , we recover the softmax

$$p(y = k \mid \psi) = \frac{e^{\psi_k}}{\sum_{k'} e^{\psi_{k'}}}$$


## The Utility Perspective

- ▶ Different priors  $\phi(\varepsilon)$  lead to different categoricals
- ▶ For  $\phi(\varepsilon) = \text{Gumbel}(\varepsilon | 0, 1)$ , we recover the softmax


$$p(y = k | \psi) = \frac{e^{\psi_k}}{\sum_{k'} e^{\psi_{k'}}}$$

- ▶ Other models: multinomial probit (Gaussian prior), multinomial logistic (logistic prior)


# The Utility Perspective



# The Utility Perspective




# The Utility Perspective



- ▶ Augment the model with *only one* error term

# The Utility Perspective



- ▶ Augment the model with *only one* error term
- ▶ Work with the joint  $p(y, \varepsilon | \psi)$
- ▶ Nice property: Amenable to stochastic optimization

## Let's Do Some Maths

- ▶ The marginal likelihood is the probability that the *realized utility*  $\psi_k + \varepsilon_k$  is greater than the others,

## Let's Do Some Maths

- ▶ The marginal likelihood is the probability that the *realized utility*  $\psi_k + \varepsilon_k$  is greater than the others,

$$p(y = k | \psi) = \text{Prob}(\psi_k + \varepsilon_k \geq \psi_{k'} + \varepsilon_{k'} \quad \forall k' \neq k)$$

## Let's Do Some Maths

- ▶ The marginal likelihood is the probability that the *realized utility*  $\psi_k + \varepsilon_k$  is greater than the others,

$$p(y = k | \psi) = \text{Prob}(\psi_k + \varepsilon_k \geq \psi_{k'} + \varepsilon_{k'} \quad \forall k' \neq k)$$

- ▶ This is an integral,

$$p(y = k | \psi) = \int_{-\infty}^{+\infty} \phi(\varepsilon_k) \left( \prod_{k' \neq k} \int_{-\infty}^{\varepsilon_k + \psi_k - \psi_{k'}} \phi(\varepsilon_{k'}) d\varepsilon_{k'} \right) d\varepsilon_k$$

## Let's Do Some Maths

- ▶ The marginal likelihood is the probability that the *realized utility*  $\psi_k + \varepsilon_k$  is greater than the others,

$$p(y = k | \psi) = \text{Prob}(\psi_k + \varepsilon_k \geq \psi_{k'} + \varepsilon_{k'} \quad \forall k' \neq k)$$

- ▶ This is an integral,

$$\begin{aligned} p(y = k | \psi) &= \int_{-\infty}^{+\infty} \phi(\varepsilon_k) \left( \prod_{k' \neq k} \int_{-\infty}^{\varepsilon_k + \psi_k - \psi_{k'}} \phi(\varepsilon_{k'}) d\varepsilon_{k'} \right) d\varepsilon_k \\ &= \int_{-\infty}^{+\infty} \phi(\varepsilon) \left( \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'}) \right) d\varepsilon \end{aligned}$$

$\Phi(\cdot)$  is the CDF of the distribution  $\phi(\cdot)$

## Let's Do Some Maths

- ▶ The marginal likelihood is the probability that the *realized utility*  $\psi_k + \varepsilon_k$  is greater than the others,

$$p(y = k | \psi) = \text{Prob}(\psi_k + \varepsilon_k \geq \psi_{k'} + \varepsilon_{k'} \quad \forall k' \neq k)$$

- ▶ This is an integral,

$$\begin{aligned} p(y = k | \psi) &= \int_{-\infty}^{+\infty} \phi(\varepsilon_k) \left( \prod_{k' \neq k} \int_{-\infty}^{\varepsilon_k + \psi_k - \psi_{k'}} \phi(\varepsilon_{k'}) d\varepsilon_{k'} \right) d\varepsilon_k \\ &= \int_{-\infty}^{+\infty} \phi(\varepsilon) \left( \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'}) \right) d\varepsilon \end{aligned}$$

$\Phi(\cdot)$  is the CDF of the distribution  $\phi(\cdot)$


- ▶ Augment the model,

$$p(y = k, \varepsilon | \psi) = \phi(\varepsilon) \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'})$$

# The Augmented Model

- We now have the augmented model,

$$p(y = k, \varepsilon | \psi) = \phi(\varepsilon) \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'})$$



# The Augmented Model

- ▶ The augmented model,

$$p(y = k, \varepsilon | \psi) = \phi(\varepsilon) \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'})$$

# The Augmented Model

- ▶ The augmented model,

$$p(y = k, \varepsilon | \psi) = \phi(\varepsilon) \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'})$$

- ▶ Nice property: The log-joint has a summation over  $k'$ ,

$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$

# The Augmented Model

- ▶ The augmented model,

$$p(y = k, \varepsilon | \psi) = \phi(\varepsilon) \prod_{k' \neq k} \Phi(\varepsilon + \psi_k - \psi_{k'})$$

- ▶ Nice property: The log-joint has a summation over  $k'$ ,

$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$

- ▶ This enables fast unbiased estimates,

1. Sample a subset of outcomes  $\mathcal{S} \subseteq \{1, \dots, K\} \setminus \{k\}$  of fixed size  $|\mathcal{S}|$
2. Compute an estimate of the log-joint in  $\mathcal{O}(|\mathcal{S}|)$  complexity

$$\log \phi(\varepsilon) + \frac{K-1}{|\mathcal{S}|} \sum_{k' \in \mathcal{S}} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$

## Augment & Reduce: Variational EM

- ▶ We are not interested in the log-joint, but in the log-marginal

## Augment & Reduce: Variational EM

- ▶ We are not interested in the log-joint, but in the log-marginal
- ▶ Variational inference relates both quantities,

$$\log p(y | \psi) \geq \mathbb{E}_{q(\varepsilon)} [\log p(y, \varepsilon | \psi) - \log q(\varepsilon)]$$

## Augment & Reduce: Variational EM

- ▶ We are not interested in the log-joint, but in the log-marginal
- ▶ Variational inference relates both quantities,

$$\log p(y | \psi) \geq \mathbb{E}_{q(\varepsilon)} [\log p(y, \varepsilon | \psi) - \log q(\varepsilon)]$$

- ▶ Maximize the bound using *variational EM*
  1. E step: Optimize w.r.t. the distribution  $q(\varepsilon)$
  2. M step: Take a gradient step w.r.t.  $\psi$  (or its parameters  $w$ )

## Example: Augment & Reduce For Classification

- ▶ Recall the classification objective,

$$\mathcal{L}_{\text{log-lik}} = \sum_n \log p(y_n | x_n, w)$$

## Example: Augment & Reduce For Classification

- ▶ Recall the classification objective,

$$\mathcal{L}_{\text{log-lik}} = \sum_n \log p(y_n | x_n, w)$$

- ▶ Replace each term with its variational bound,

$$\mathcal{L}_{\text{bound}} = \sum_n \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \log p(y_n, \varepsilon^{(n)} | x_n, w) - \log q(\varepsilon^{(n)}) \right]$$

## Example: Augment & Reduce For Classification

- ▶ Recall the classification objective,

$$\mathcal{L}_{\text{log-lik}} = \sum_n \log p(y_n | x_n, w)$$

- ▶ Replace each term with its variational bound,

$$\mathcal{L}_{\text{bound}} = \sum_n \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \log p(y_n, \varepsilon^{(n)} | x_n, w) - \log q(\varepsilon^{(n)}) \right]$$


- ▶ Algorithm

1. Subsample datapoints  $\mathcal{B} \subseteq \{1, \dots, N\}$
2. For each  $n \in \mathcal{B}$ , subsample classes  $\mathcal{S} \subseteq \{1, \dots, K\} \setminus \{y_n\}$
3. (E step) For each  $n \in \mathcal{B}$ , update its  $q(\varepsilon^{(n)})$   $\mathcal{O}(|\mathcal{S}|)$
4. (M step) For each  $n \in \mathcal{B}$ , compute gradient w.r.t.  $w$   $\mathcal{O}(|\mathcal{S}|)$
5. (M step) Take gradient step for  $w$
6. Repeat

## Example: Augment & Reduce For Classification

- ▶ Recall the log-joint in the augmented model,


$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$



## Example: Augment & Reduce For Classification

- Recall the log-joint in the augmented model,

$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$




- Consider the gradient of the bound in the M step,

$$\nabla_w \mathcal{L}_{\text{bound}} = \nabla_w \sum_n \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \log p(y_n, \varepsilon^{(n)} | x_n, w) - \log q(\varepsilon^{(n)}) \right]$$

## Example: Augment & Reduce For Classification

- Recall the log-joint in the augmented model,

$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$




- Consider the gradient of the bound in the M step,

$$\begin{aligned}\nabla_w \mathcal{L}_{\text{bound}} &= \nabla_w \sum_n \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \log p(y_n, \varepsilon^{(n)} | x_n, w) - \log q(\varepsilon^{(n)}) \right] \\ &= \sum_n \sum_{k' \neq y_n} \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \nabla_w \log \Phi(\varepsilon^{(n)} + w_{y_n}^\top x_n - w_{k'}^\top x_n) \right]\end{aligned}$$

## Example: Augment & Reduce For Classification

- Recall the log-joint in the augmented model,

$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$




- Consider the gradient of the bound in the M step,

$$\begin{aligned}\nabla_w \mathcal{L}_{\text{bound}} &= \nabla_w \sum_n \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \log p(y_n, \varepsilon^{(n)} | x_n, w) - \log q(\varepsilon^{(n)}) \right] \\ &= \sum_n \sum_{k' \neq y_n} \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \nabla_w \log \Phi(\varepsilon^{(n)} + w_{y_n}^\top x_n - w_{k'}^\top x_n) \right] \\ &\approx \frac{N}{|\mathcal{B}|} \frac{K-1}{|\mathcal{S}|} \sum_{n \in \mathcal{B}} \sum_{k' \in \mathcal{S}_n} \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \nabla_w \log \Phi(\varepsilon^{(n)} + w_{y_n}^\top x_n - w_{k'}^\top x_n) \right]\end{aligned}$$

## Example: Augment & Reduce For Classification

- Recall the log-joint in the augmented model,

$$\log p(y = k, \varepsilon | \psi) = \log \phi(\varepsilon) + \sum_{k' \neq k} \log \Phi(\varepsilon + \psi_k - \psi_{k'})$$



- Consider the gradient of the bound in the M step,

$$\begin{aligned}\nabla_w \mathcal{L}_{\text{bound}} &= \nabla_w \sum_n \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \log p(y_n, \varepsilon^{(n)} | x_n, w) - \log q(\varepsilon^{(n)}) \right] \\ &= \sum_n \sum_{k' \neq y_n} \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \nabla_w \log \Phi(\varepsilon^{(n)} + w_{y_n}^\top x_n - w_{k'}^\top x_n) \right] \\ &\approx \frac{N}{|\mathcal{B}|} \frac{K-1}{|\mathcal{S}|} \sum_{n \in \mathcal{B}} \sum_{k' \in \mathcal{S}_n} \mathbb{E}_{q(\varepsilon^{(n)})} \left[ \nabla_w \log \Phi(\varepsilon^{(n)} + w_{y_n}^\top x_n - w_{k'}^\top x_n) \right]\end{aligned}$$

## Augment & Reduce For Softmax Model

- ▶ The softmax model is special
  - ▶ We can compute the optimal  $q(\varepsilon)$  distribution
  - ▶ We can compute the integrals

## Augment & Reduce For Softmax Model

- ▶ The softmax model is special
  - ▶ We can compute the optimal  $q(\varepsilon)$  distribution
  - ▶ We can compute the integrals
- ▶ The optimal variational distribution is Gumbel,

$$q^*(\varepsilon) = \text{Gumbel}(\log \eta^*, 1), \quad \eta^* = 1 + \sum_{k' \neq k} e^{\psi_{k'} - \psi_k}$$

## Augment & Reduce For Softmax Model

- ▶ The softmax model is special
  - ▶ We can compute the optimal  $q(\varepsilon)$  distribution
  - ▶ We can compute the integrals
- ▶ The optimal variational distribution is Gumbel,

$$q^*(\varepsilon) = \text{Gumbel}(\log \eta^*, 1), \quad \eta^* = 1 + \sum_{k' \neq k} e^{\psi_{k'} - \psi_k}$$

- ▶ Instead, set

$$q^*(\varepsilon) = \text{Gumbel}(\log \eta, 1)$$

# Augment & Reduce For Softmax Model

- ▶ The softmax model is special
  - ▶ We can compute the optimal  $q(\varepsilon)$  distribution
  - ▶ We can compute the integrals
- ▶ The optimal variational distribution is Gumbel,

$$q^*(\varepsilon) = \text{Gumbel}(\log \eta^*, 1), \quad \eta^* = 1 + \sum_{k' \neq k} e^{\psi_{k'} - \psi_k}$$

- ▶ Instead, set

$$q^*(\varepsilon) = \text{Gumbel}(\log \eta, 1)$$

- ▶ Estimate the optimal natural parameter,

$$\tilde{\eta} = 1 + \frac{K-1}{|S|} \sum_{k' \in S} e^{\psi_{k'} - \psi_k}$$

(to update  $\eta$ , take a step in the direction of the natural gradient)

## Augment & Reduce For Other Models

- ▶ For other models, the expectations are intractable

## Augment & Reduce For Other Models

- ▶ For other models, the expectations are intractable
- ▶ We form Monte Carlo gradient estimators using the reparameterization trick

## Augment & Reduce For Other Models

- ▶ For other models, the expectations are intractable
- ▶ We form Monte Carlo gradient estimators using the reparameterization trick
- ▶ Useful for both E and M steps

# Experiments

- ▶ Experiments: Linear classification
- ▶ Maximum likelihood estimation
- ▶ 5 datasets

dataset	$N_{\text{train}}$	$N_{\text{test}}$	covariates	classes	minibatch (obs.)	minibatch (classes)	iterations
MNIST	60,000	10,000	784	10	500	1	35,000
Bibtex	4,880	2,413	1,836	148	488	20	5,000
Omniglot	25,968	6,492	784	1,623	541	50	45,000
EURLex-4K	15,539	3,809	5,000	896	279	50	100,000
AmazonCat-13K	1,186,239	306,782	203,882	2,919	1,987	60	5,970

# Experiments

- ▶ Comparisons against:
  - ▶ Exact softmax (only for MNIST and Bibtex)

# Experiments

- ▶ Comparisons against:
  - ▶ Exact softmax (only for MNIST and Bibtex)
  - ▶ One-vs-each (OVE) bound,

$$\mathcal{L}_{\text{OVE}} = \sum_{k' \neq k} \log \sigma(\psi_k - \psi_{k'})$$

(it is a bound on the softmax)


# Experiments

## ► Time complexity

dataset	OVE ( <a href="#">Titlias, 2016</a> )	A&R [this paper]		
		softmax	multi. probit	multi. logistic
MNIST	0.336 s	0.337 s	0.431 s	0.511 s
Bibtex	0.181 s	0.188 s	0.244 s	0.246 s
Omniglot	4.47 s	4.65 s	5.63 s	5.57 s
EURLex-4K	5.54 s	5.65 s	6.46 s	6.23 s
AmazonCat-13K	2.80 h	2.80 h	2.82 h	2.91 h

# Experiments

## ► Quality of the bound



# Experiments

- ▶ Quality of the classification weights  $w_k$  (predictive performance)

dataset	exact		softmax model OVE (Titsias, 2016)		A&R [this paper]		multi. probit		multi. logistic	
	log lik	acc	log lik	acc	log lik	acc	log lik	acc	log lik	acc
MNIST	-0.261	0.927	-0.276	0.919	<b>-0.271</b>	<b>0.924</b>	-0.302	0.918	-0.287	0.917
Bibtex	-3.188	0.361	-3.300	0.352	<b>-3.036</b>	<b>0.361</b>	-4.184	0.346	-3.151	0.353
Omniglot	—	—	-5.667	0.179	<b>-5.171</b>	<b>0.201</b>	-7.350	0.178	-5.395	0.184
EURLex-4K	—	—	<b>-4.241</b>	<b>0.247</b>	-4.593	0.207	-4.193	0.263	-4.299	0.226
AmazonCat-13K	—	—	-3.880	0.388	<b>-3.795</b>	<b>0.420</b>	-3.593	0.411	-4.081	0.350

# Conclusion

- ▶ A method to scale up training for models involving large categorical distributions
- ▶ Stochastic variational EM
- ▶ Controlled complexity ( $|\mathcal{S}|$  is a parameter)
- ▶ Can be embedded in many different models
- ▶ Not limited to maximum likelihood estimation

Thank you for your attention!

